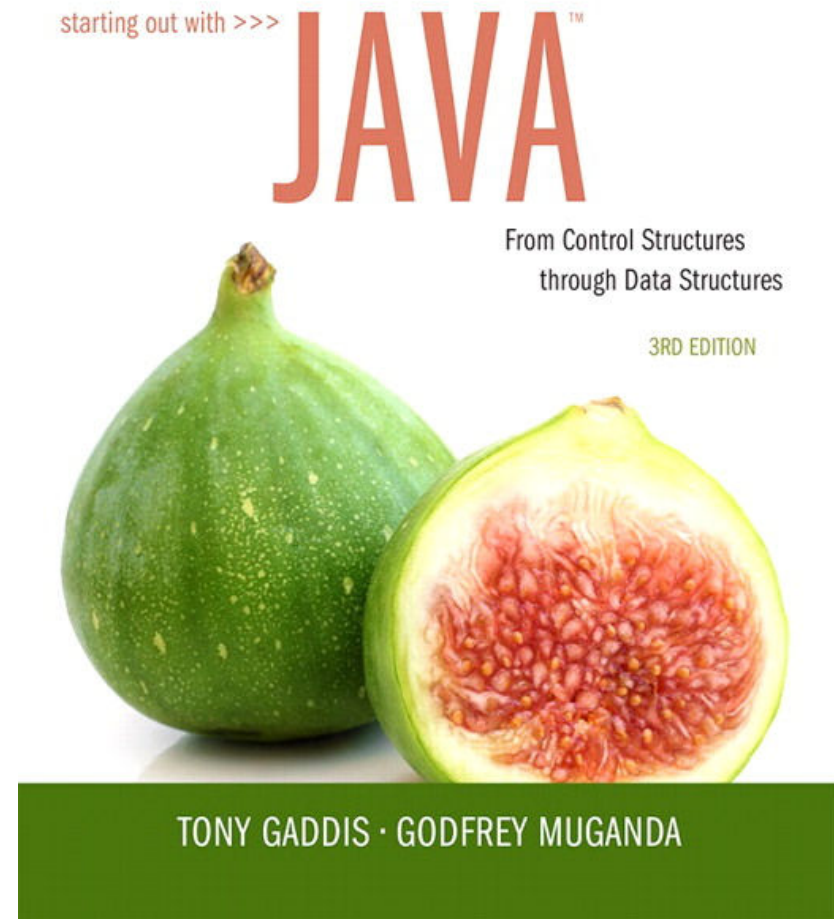


CHAPTER 11

Exceptions and Advanced File I/O



Chapter Topics

Chapter 11 discusses the following main topics:

- Handling Exceptions
- Throwing Exceptions
- More about Input/Output Streams
- Advanced Topics:
 - Binary Files,
 - Random Access Files, and
 - Object Serialization

Handling Exceptions

- An exception is an object that is generated as the result of an error or an unexpected event.
- Exception are said to have been “thrown.”
- It is the programmers responsibility to write code that detects and handles exceptions.
- Unhandled exceptions will crash a program.
- Example: [BadArray.java](#)
- Java allows you to create exception handlers.

Handling Exceptions

- An *exception handler* is a section of code that gracefully responds to exceptions.
- The process of intercepting and responding to exceptions is called *exception handling*.
- The *default exception handler* deals with unhandled exceptions.
- The default exception handler prints an error message and crashes the program.

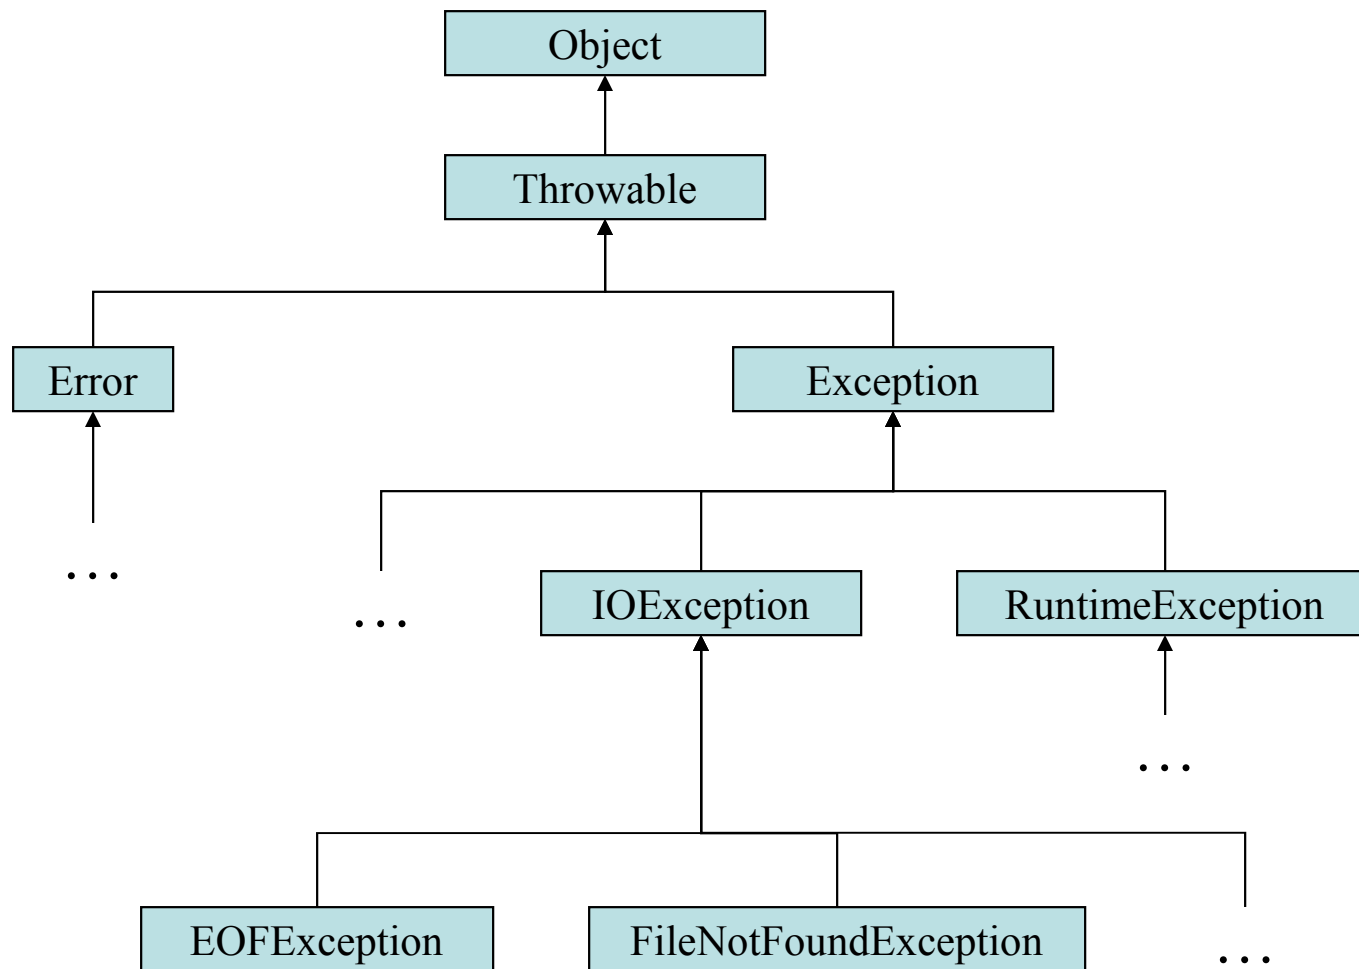
Exception Classes

- An exception is an object.
- Exception objects are created from classes in the Java API hierarchy of exception classes.
- All of the exception classes in the hierarchy are derived from the `Throwable` class.
- `Error` and `Exception` are derived from the `Throwable` class.

Exception Classes

- Classes that are derived from `Error`:
 - are for exceptions that are thrown when critical errors occur. (i.e.)
 - an internal error in the Java Virtual Machine, or
 - running out of memory.
- Applications should not try to handle these errors because they are the result of a serious condition.
- Programmers should handle the exceptions that are instances of classes that are derived from the `Exception` class.

Exception Classes



Handling Exceptions

- To handle an exception, you use a *try* statement.

```
try
{
    (try block statements...)
}
catch (ExceptionType ParameterName)
{
    (catch block statements...)
}
```

- First the keyword `try` indicates a block of code will be attempted (the curly braces are required).
- This block of code is known as a *try block*.

Handling Exceptions

- A *try block* is:
 - one or more statements that are executed, and
 - can potentially throw an exception.
- The application will not halt if the try block throws an exception.
- After the try block, a `catch` clause appears.

Handling Exceptions

- A catch clause begins with the key word `catch`:

`catch (ExceptionType ParameterName)`

- *ExceptionType* is the name of an exception class and
 - *ParameterName* is a variable name which will reference the exception object if the code in the try block throws an exception.
- The code that immediately follows the catch clause is known as a *catch block* (the curly braces are required).
- The code in the catch block is executed if the try block throws an exception.

Handling Exceptions

- This code is designed to handle a `FileNotFoundException` if it is thrown.

```
try
{
    File file = new File ("MyFile.txt");
    Scanner inputFile = new Scanner(file);
}
catch (FileNotFoundException e)
{
    System.out.println("File not found.");
}
```

- The Java Virtual Machine searches for a `catch` clause that can deal with the exception.
- Example: [OpenFile.java](#)

Handling Exceptions

- The parameter must be of a type that is compatible with the thrown exception's type.
- After an exception, the program will continue execution at the point just past the catch block.

Handling Exceptions

- Each exception object has a method named `getMessage` that can be used to retrieve the default error message for the exception.
- Example:
 - [ExceptionMessage.java](#)
 - [ParseIntError.java](#)

Polymorphic References To Exceptions

- When handling exceptions, you can use a polymorphic reference as a parameter in the `catch` clause.
- Most exceptions are derived from the `Exception` class.
- A `catch` clause that uses a parameter variable of the `Exception` type is capable of catching any exception that is derived from the `Exception` class.

Polymorphic References To Exceptions

```
try
{
    number = Integer.parseInt(str);
}
catch (Exception e)
{
    System.out.println("The following error occurred: "
                        + e.getMessage());
}
```

- The Integer class's parseInt method throws a NumberFormatException object.
- The NumberFormatException class is derived from the Exception class.

Handling Multiple Exceptions

- The code in the try block may be capable of throwing more than one type of exception.
- A `catch` clause needs to be written for each type of exception that could potentially be thrown.
- The JVM will run the first compatible `catch` clause found.
- The `catch` clauses must be listed from most specific to most general.
- Example: [SalesReport.java](#), [SalesReport2.java](#)

Exception Handlers

- There can be many polymorphic catch clauses.
- A try statement may have only one catch clause for each specific type of exception.

```
try
{
    number = Integer.parseInt(str);
}
catch (NumberFormatException e)
{
    System.out.println("Bad number format.");
}
catch (NumberFormatException e) // ERROR!!!
{
    System.out.println(str + " is not a number.");
}
```

Exception Handlers

- The `NumberFormatException` class is derived from the `IllegalArgumentException` class.

```
try
{
    number = Integer.parseInt(str);
}
catch (IllegalArgumentException e)
{
    System.out.println("Bad number format.");
}
catch (NumberFormatException e) // ERROR!!!
{
    System.out.println(str + " is not a number.");
}
```

Exception Handlers

- The previous code could be rewritten to work, as follows, with no errors:

```
try
{
    number = Integer.parseInt(str);
}
catch (NumberFormatException e)
{
    System.out.println(str +
                        " is not a number.");
}
catch (IllegalArgumentException e) //OK
{
    System.out.println("Bad number format.");
}
```

The `finally` Clause

- The try statement may have an optional `finally` clause.
- If present, the `finally` clause must appear after all of the catch clauses.

```
try
{
    (try block statements...)
}
catch (ExceptionType ParameterName)
{
    (catch block statements...)
}
finally
{
    (finally block statements...)
}
```

The *finally* Clause

- The *finally block* is one or more statements,
 - that are always executed after the try block has executed and
 - after any catch blocks have executed if an exception was thrown.
- The statements in the finally block execute whether an exception occurs or not.


The Stack Trace

- The *call stack* is an internal list of all the methods that are currently executing.
- A *stack trace* is a list of all the methods in the call stack.
- It indicates:
 - the method that was executing when an exception occurred and
 - all of the methods that were called in order to execute that method.
- Example: [StackTrace.java](#)

Multi-Catch (Java 7)

- Beginning in Java 7, you can specify more than one exception in a catch clause:

```
try
{
}
catch (NumberFormatException | InputMismatchException ex)
{
}
```



Separate the exceptions with
the | character.

Uncaught Exceptions

- When an exception is thrown, it cannot be ignored.
- It must be handled by the program, or by the default exception handler.
- When the code in a method throws an exception:
 - normal execution of that method stops, and
 - the JVM searches for a compatible exception handler inside the method.

Uncaught Exceptions

- If there is no exception handler inside the method:
 - control of the program is passed to the previous method in the call stack.
 - If that method has no exception handler, then control is passed again, up the call stack, to the previous method.
- If control reaches the `main` method:
 - the main method must either handle the exception, or
 - the program is halted and the default exception handler handles the exception.

The Catch or Declare Rule

- Most ordinary exceptions that might be thrown within a method must be accounted for in one of two ways:
 1. The code that can throw an exception is placed within a **try** block, and the possible exception is caught in a **catch** block within the same method
 2. The possible exception can be declared at the start of the method definition by placing the exception class name in a **throws** clause

The Catch or Declare Rule

- The first technique handles an exception in a **catch** block
- The second technique is a way to shift the exception handling responsibility to the method that invoked the exception throwing method
- The invoking method must handle the exception, unless it too uses the same technique to "pass the buck"
- Ultimately, every exception that is thrown should eventually be caught by a **catch** block in some method that does not just declare the exception class in a **throws** clause

The Catch or Declare Rule

- In any one method, both techniques can be mixed
 - Some exceptions may be caught, and others may be declared in a **throws** clause
- However, these techniques must be used consistently with a given exception
 - If an exception is not declared, then it must be handled within the method
 - If an exception is declared, then the responsibility for handling it is shifted to some other calling method
 - Note that if a method definition encloses an invocation of a second method, and the second method can throw an exception and does not catch it, then the first method must catch or declare it

Checked and Unchecked Exceptions

- Exceptions that are subject to the catch or declare rule are called *checked* exceptions
 - The compiler checks to see if they are accounted for with either a catch block or a throws clause
 - The classes **Throwable**, **Exception**, and all descendants of the class **Exception** are checked exceptions
- All other exceptions are *unchecked* exceptions
- The class **Error** and all its descendant classes are called *error classes*
 - Error classes are *not* subject to the Catch or Declare Rule

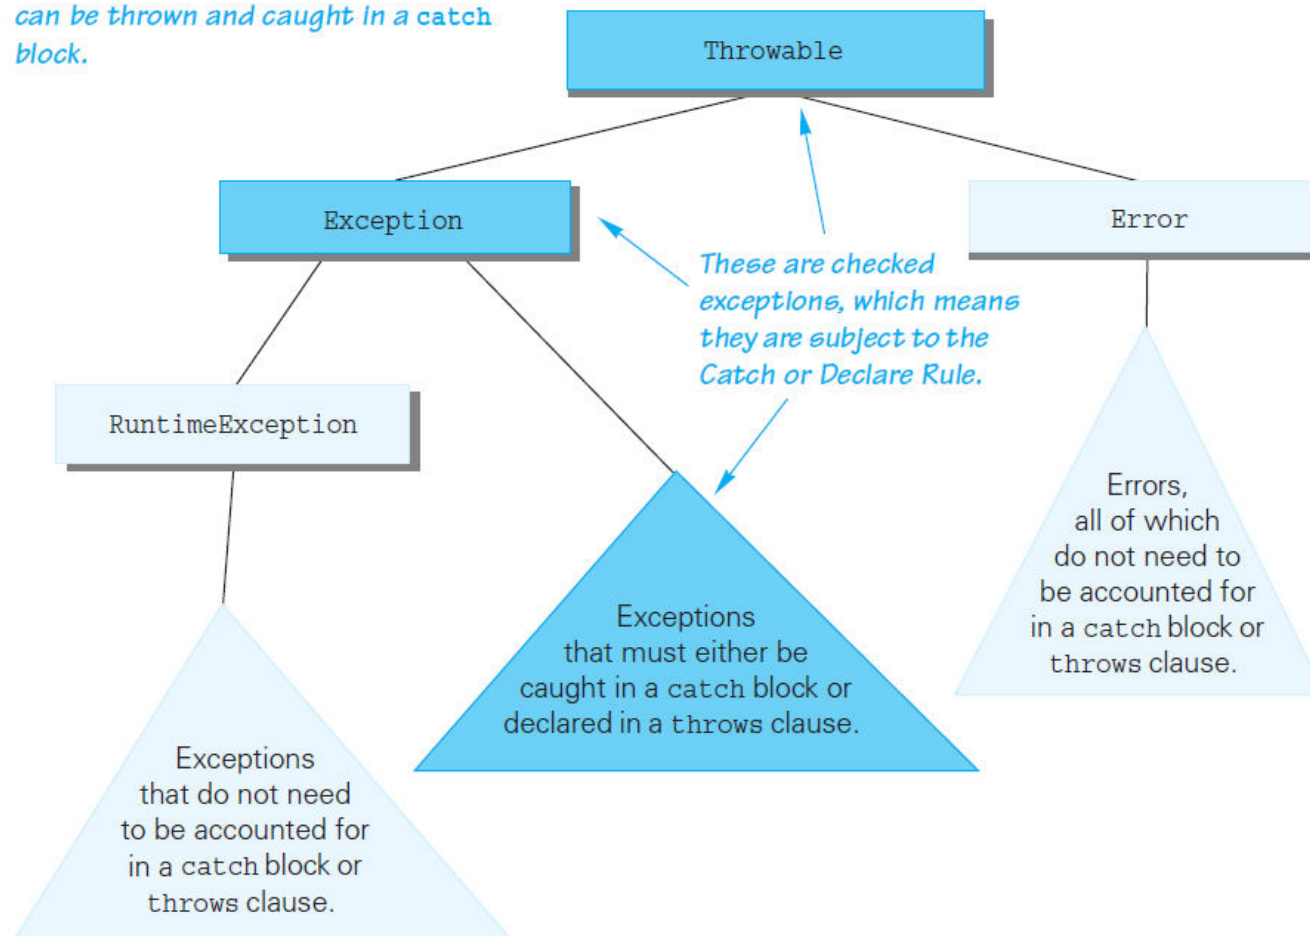
Exceptions to the Catch or Declare Rule

- Checked exceptions must follow the Catch or Declare Rule
 - Programs in which these exceptions can be thrown will not compile until they are handled properly
- Unchecked exceptions are exempt from the Catch or Declare Rule
 - Programs in which these exceptions are thrown simply need to be corrected, as they result from some sort of error

Hierarchy of Throwable Objects

Display 9.11 Hierarchy of Throwable Objects

All descendents of the class Throwable can be thrown and caught in a catch block.



Throwing Exceptions

- You can write code that:
 - throws one of the standard Java exceptions, or
 - an instance of a custom exception class that you have designed.
- The `throw` statement is used to manually throw an exception.

```
throw new ExceptionType(MessageString) ;
```

- The `throw` statement causes an exception object to be created and thrown.

Throwing Exceptions

- The *MessageString* argument contains a custom error message that can be retrieved from the exception object's `getMessage` method.
- If you do not pass a message to the constructor, the exception will have a null message.

```
throw new Exception("Out of fuel");
```

– *Note: Don't confuse the `throw` statement with the `throws` clause.*

- Example: [DieExceptionDemo.java](#)

Creating Exception Classes

- You can create your own exception classes by deriving them from the `Exception` class or one of its derived classes.
- Example:
 - [BankAccount.java](#)
 - [NegativeStartingBalance.java](#)
 - [AccountTest.java](#)

Creating Exception Classes

- Some examples of exceptions that can affect a bank account:
 - A negative starting balance is passed to the constructor.
 - A negative interest rate is passed to the constructor.
 - A negative number is passed to the deposit method.
 - A negative number is passed to the withdraw method.
 - The amount passed to the withdraw method exceeds the account's balance.
- We can create exceptions that represent each of these error conditions.

@exception Tag in Documentation Comments

- General format

`@exception` *ExceptionName Description*

- The following rules apply

- The `@exception` tag in a method's documentation comment must appear after the general description of the method.
- The description can span several lines. It ends at the end of the documentation comment (the `*/` symbol) or at the beginning of another tag.

Binary Files

Binary Files

- The way data is stored in memory is sometimes called the *raw binary format*.
- Data can be stored in a file in its raw binary format.
- A file that contains binary data is often called a *binary file*.
- Storing data in its binary format is more efficient than storing it as text.
- There are some types of data that should only be stored in its raw binary format.

Binary Files

- Binary files cannot be opened in a text editor such as Notepad.
- To write data to a binary file you must create objects from the following classes:
 - **FileOutputStream** - allows you to open a file for writing binary data. It provides only basic functionality for writing bytes to the file.
 - **DataOutputStream** - allows you to write data of any primitive type or String objects to a binary file. Cannot directly access a file. It is used in conjunction with a `FileOutputStream` object that has a connection to a file.

Binary Files

- A `DataOutputStream` object is wrapped around a `FileOutputStream` object to write data to a binary file.

```
FileOutputStream fstream = new  
    FileOutputStream("MyInfo.dat");  
DataOutputStream outputFile = new  
    DataOutputStream(fstream);
```

- If the file that you are opening with the `FileOutputStream` object already exists, it will be erased and an empty file by the same name will be created.

Binary Files

- These statements can be combined into one.

```
DataOutputStream outputFile = new  
    OutputStream(new  
        FileOutputStream("MyInfo.dat")) ;
```

- Once the `DataOutputStream` object has been created, you can use it to write binary data to the file.
- Example: [WriteBinaryFile.java](#)

Binary Files

- To open a binary file for input, you wrap a `DataInputStream` object around a `FileInputStream` object.

```
FileInputStream fstream = new
    FileInputStream("MyInfo.dat");
DataInputStream inputFile = new
    DataInputStream(fstream);
```

- These two statements can be combined into one.

```
DataInputStream inputFile = new
    DataInputStream(new
        FileInputStream("MyInfo.dat"));
```

Binary Files

- The `FileInputStream` constructor will throw a `FileNotFoundException` if the file named by the string argument cannot be found.
- Once the `DataInputStream` object has been created, you can use it to read binary data from the file.
- Example:
 - [ReadBinaryFile.java](#)

Writing and Reading Strings

- To write a string to a binary file, use the `DataOutputStream` class's `writeUTF` method.
- This method writes its `String` argument in a format known as *UTF-8 encoding*.
 - Just before writing the string, this method writes a two-byte integer indicating the number of bytes that the string occupies.
 - Then, it writes the string's characters in Unicode. (UTF stands for Unicode Text Format.)
- The `DataInputStream` class's `readUTF` method reads from the file.

Writing and Reading Strings

- To write a string to a file:
`String name = "Chloe";`
`outputFile.writeUTF(name);`
- To read a string from a file:
`String name = inputFile.readUTF();`
- The `readUTF` method will correctly read a string only when the string was written with the `writeUTF` method.
- Example:
 - [WriteUTF.java](#)
 - [ReadUTF.java](#)

Appending Data to Binary Files

- The `FileOutputStream` constructor takes an optional second argument which must be a `boolean` value.
- If the argument is `true`, the file will not be erased if it exists; new data will be written to the end of the file.
- If the argument is `false`, the file will be erased if it already exists.

```
FileOutputStream fstream = new  
    FileOutputStream("MyInfo.dat", true);  
DataOutputStream outputFile = new  
    DataOutputStream(fstream);
```

Random Access Files

- Text files and the binary files previously shown use *sequential file access*.
- With sequential access:
 - The first time data is read from the file, the data will be read from its beginning.
 - As the reading continues, the file's read position advances sequentially through the file's contents.
- Sequential file access is useful in many circumstances.
- If the file is very large, locating data buried deep inside it can take a long time.

Random Access Files

- Java allows a program to perform *random file access*.
- In random file access, a program may immediately jump to any location in the file.
- To create and work with random access files in Java, you use the `RandomAccessFile` class.

`RandomAccessFile(String filename, String mode)`

- *filename*: the name of the file.
- *mode*: a string indicating the mode in which you wish to use the file.
 - "r" = reading
 - "rw" = for reading and writing.

Random Access Files

```
// Open a file for random reading.  
RandomAccessFile randomFile = new  
    RandomAccessFile("MyData.dat", "r");  
// Open a file for random reading and writing.  
RandomAccessFile randomFile = new  
    RandomAccessFile("MyData.dat", "rw");
```

- When opening a file in "r" mode where the file does not exist, a `FileNotFoundException` will be thrown.
- Opening a file in "r" mode and trying to write to it will throw an `IOException`.
- If you open an existing file in "rw" mode, it will not be deleted and the file's existing content will be preserved.

Random Access Files

- Items in a sequential access file are accessed one after the other.
- Items in a random access file are accessed in any order.
- If you open a file in "rw" mode and the file does not exist, it will be created.
- A file that is opened or created with the `RandomAccessFile` class is treated as a binary file.

Random Access Files

- The `RandomAccessFile` class has:
 - the same methods as the `DataOutputStream` class for writing data, and
 - the same methods as the `DataInputStream` class for reading data.
- The `RandomAccessFile` class can be used to sequentially process a binary file.
- Example: [WriteLetters.java](#)

The File Pointer

- The `RandomAccessFile` class treats a file as a stream of bytes.
- The bytes are numbered:
 - the first byte is byte 0.
 - The last byte's number is one less than the number of bytes in the file.
- These byte numbers are similar to an array's subscripts, and are used to identify locations in the file.
- Internally, the `RandomAccessFile` class keeps a long integer value known as the *file pointer*.

The File Pointer

- The *file pointer* holds the byte number of a location in the file.
- When a file is first opened, the file pointer is set to 0.
- When an item is read from the file, it is read from the byte that the file pointer points to.
- Reading also causes the file pointer to advance to the byte just beyond the item that was read.
- If another item is immediately read, the reading will begin at that point in the file.

The File Pointer

- An `EOFException` is thrown when a read causes the file pointer to go beyond the size of the file.
- Writing also takes place at the location pointed to by the file pointer.
- If the file pointer points to the end of the file, data will be written to the end of the file.
- If the file pointer holds the number of a byte within the file, at a location where data is already stored, a write will overwrite the data at that point.

The File Pointer

- The `RandomAccessFile` class lets you move the file pointer.
- This allows data to be read and written at any byte location in the file.
- The `seek` method is used to move the file pointer.

`rndFile.seek(long position) ;`

- The argument is the number of the byte that you want to move the file pointer to.

The File Pointer

```
RandomAccessFile file = new  
    RandomAccessFile("MyInfo.dat", "r");  
file.seek(99);  
byte b = file.readByte();
```

- Example: [ReadRandomLetters.java](#)

Object Serialization

- If an object contains other types of objects as fields, saving its contents can be complicated.
- Java allows you to *serialize* objects, which is a simpler way of saving objects to a file.
- When an object is serialized, it is converted into a series of bytes that contain the object's data.
- If the object is set up properly, even the other objects that it might contain as fields are automatically serialized.
- The resulting set of bytes can be saved to a file for later retrieval.

Object Serialization

- For an object to be serialized, its class must implement the `Serializable` interface.
- The `Serializable` interface has no methods or fields.
- It is used only to let the Java compiler know that objects of the class might be serialized.
- If a class contains objects of other classes as fields, those classes must also implement the `Serializable` interface, in order to be serialized.
- Example: [BankAccount2.java](#)

Object Serialization

- The `String` class, as many others in the Java API, implements the `Serializable` interface.
- To write a serialized object to a file, you use an `ObjectOutputStream` object.
- The `ObjectOutputStream` class is designed to perform the serialization process.
- To write the bytes to a file, an output stream object is needed.

```
FileOutputStream outputStream = new  
    FileOutputStream("Objects.dat");  
ObjectOutputStream objectOutputStream = new  
    ObjectOutputStream(outputStream);
```

Object Serialization

- To serialize an object and write it to the file, the `ObjectOutputStream` class's `writeObject` method is used.

```
BankAccount2 account = new  
    BankAccount(25000.0) ;  
objectOutputFile.writeObject(account) ;
```

- The `writeObject` method throws an `IOException` if an error occurs.
- The process of reading a serialized object's bytes and constructing an object from them is known as *deserialization*.

Object Serialization

- To deserialize an object an `ObjectInputStream` object is used in conjunction with a `FileInputStream` object.

```
FileInputStream inStream = new  
    FileInputStream("Objects.dat");  
ObjectInputStream objectInputFile = new  
    ObjectInputStream(inStream);
```

- To read a serialized object from the file, the `ObjectInputStream` class's `readObject` method is used.

```
BankAccount2 account;  
account = (BankAccount2)  
    objectInputFile.readObject();
```

Object Serialization

- The `readObject` method returns the deserialized object.
 - *Notice that you must cast the return value to the desired class type.*
- The `readObject` method throws a number of different exceptions if an error occurs.
- Examples:
 - [SerializeObjects.java](#)
 - [DeserializeObjects.java](#)