

Chapter 17: Sorting, Searching, and Algorithm Analysis

**Starting Out with Java
From Control Structures through Data Structures**

by Tony Gaddis and Godfrey Muganda



Copyright © 2007 Pearson Education, Inc. Publishing as Pearson Addison-Wesley

Chapter Topics

- Introduction to Sorting Algorithms
- Introduction to Searching Algorithms
- Analysis of Algorithms

What is Sorting?

- An array $A[]$ of N numbers is **sorted in ascending order** if the array entries increase (or never decrease) as indices increase:

$$A[0] \leq A[1] \leq \dots \leq A[N-2] \leq A[N-1]$$

What is Sorting?

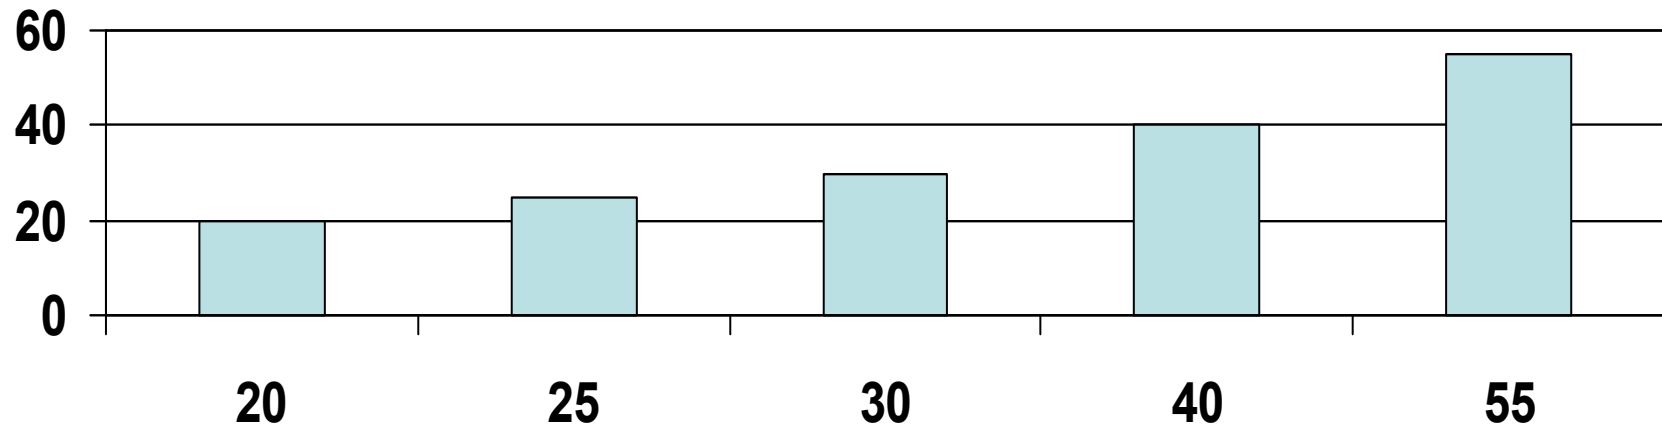
Example of an array sorted in ascending order:

20 25 30 40 55

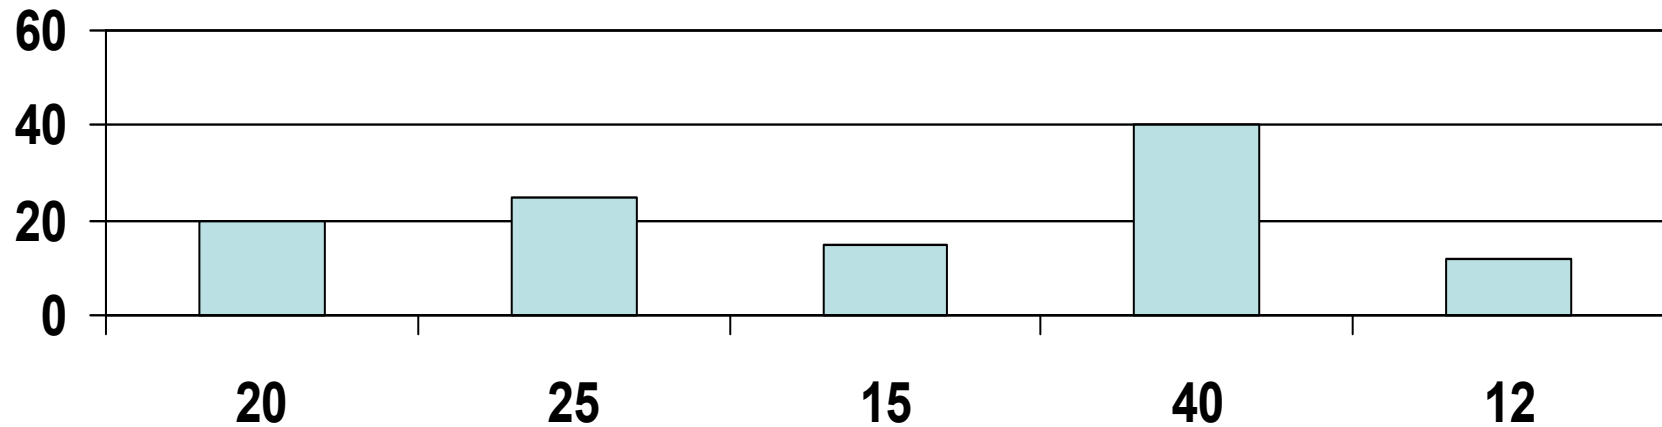
Example of an array that is not sorted in ascending order:

20 25 15 40 12

Example of a Sorted Array



An Array That is Not Sorted



Sorting Algorithms

There are many strategies for sorting arrays.

Among them:

- Bubble sort
- Selection sort
- Insertion sort
- Quicksort

Notation

- Let $A[]$ be an array of length N
- Let last be an index in the range of the array:

$$0 \leq \text{last} \leq N - 1$$

- $A[0..\text{last}]$ denotes the portion of the array consisting of $A[0]$, $A[1]$, ..., $A[\text{last}]$

A Simple Sorting Strategy

```
for (last =  $N - 1$ ; last  $\geq 1$ ; last --)
{
    Move the largest entry in  $A[0 \dots \text{last}]$  to  $A[\text{last}]$ 
}
```

A Simple Sorting Strategy

How it works on 15 35 20 10 25.

Portion of array already sorted in orange.

Largest value in unsorted portion of array in bright blue.

- 15 35 20 10 25
- Move largest of A[0..4] to A[4]: 15 20 10 25 35
- Move largest of A[0..3] to A[3]: 15 20 10 25 35
- Move largest of A[0..2] to A[2]: 15 10 20 25 35
- Move largest of A[0..1] to A[1]: 10 15 20 25 35

```
for (last = 4; last >= 1; last--)  
{  
    Move largest of A[0..last] to A[last]  
}
```

Selection and Bubble Sort

- Selection and Bubble Sort are similar: both use the Simple Sorting Strategy of the previous slide.
- Selection and Bubble Sort differ in how they implement the step:
Move the largest of $A[0..last]$ to $A[last]$

Bubble Sort

Bubble Sort moves the largest entry to the end of $A[0..last]$ by comparing and swapping **adjacent** elements as an index sweeps through the unsorted portion of the array:

15 35 20 10 25	//Compare $A[0]$, $A[1]$, no swap
15 35 20 10 25	//Compare $A[1]$, $A[2]$, swap
15 20 35 10 25	
15 20 35 10 25	//Compare $A[2]$, $A[3]$, swap
15 20 10 35 25	
15 20 10 35 25	//Compare $A[3]$, $A[4]$, swap
15 20 10 25 35	//Largest is at $A[4]$

Bubble sort uses an index to keep track of which pair of adjacent elements should be swapped during a sweep through A[0..last].

15 35 20 10 25	// index = 0, Compare A[0], A[1]
15 35 20 10 25	// index = 1, Compare A[1], A[2], swap
15 20 35 10 25	
15 20 35 10 25	// index = 2, Compare A[2], A[3], swap
15 20 10 35 25	
15 20 10 35 25	//index = 3, Compare A[3], A[4], swap
15 20 10 25 35	//Largest is at A[4]

Bubble Sort

To accomplish the step:

move largest of $A[0..last]$ to $A[last]$

Bubble Sort sweeps an $index$ from 0 to $last-1$, swapping adjacent entries to put the largest element seen so far at $index + 1$:

```
for (index = 0; index <= last - 1; index++)  
{ //swap adjacent elements if necessary  
  if (A[index] > A[index+1])  
  {  
    int temp = A[index];  
    A[index] = A[index+1];  
    A[index + 1] = temp;  
  }  
}
```

Bubble Sort

To sort an array $A[0..N-1]$:

```
for (int last = N - 1; last >= 1; last --)
{
    // Move the largest entry in  $A[0...last]$  to  $A[last]$ 
    for (int index = 0; index <= last-1; index++)
    {
        //swap adjacent elements if necessary
        if (A[index] > A[index+1])
        {
            int temp = A[index];
            A[index] = A[index+1];
            A[index + 1] = temp;
        }
    }
}
```

Selection Sort

Selection sort, like Bubble sort, is based on a strategy that repeatedly executes the step:

Move the largest of $A[0..last]$ to $A[last]$

Selection Sort

Selection Sort moves the largest entry to the end of $A[0..last]$ by determining the position $maxIndex$ of the largest entry, and then swapping $A[maxIndex]$ with $A[last]$:

Selection Sort uses 2 variables:

1. $index$: keeps track of the portion $A[0..index]$ that has already been examined: $index$ will range from 0 to $last$.
2. $maxIndex$: keeps track of the position of the largest entry in $A[0..index]$. When $index$ equals $last$, the variable $maxIndex$ will be the position of the largest entry in $A[0..last]$.

Determining the position of the largest entry

The portion of $A[0..\text{last}]$ already examined while looking for the largest entry is in orange.

The largest entry found in the portion already examined is in bright blue.

20 35 25 15 50 40 60 45 30	index = 0, maxIndex = 0
20 35 25 15 50 40 60 45 30	index = 1, maxIndex = 1
20 35 25 15 50 40 60 45 30	index = 2, maxIndex = 1
20 35 25 15 50 40 60 45 30	index = 3, maxIndex = 1
20 35 25 15 50 40 60 45 30	index = 4, maxIndex = 4
20 35 25 15 50 40 60 45 30	index = 5, maxIndex = 4
20 35 25 15 50 40 60 45 30	index = 6, maxIndex = 6
20 35 25 15 50 40 60 45 30	index = 7, maxIndex = 6
20 35 25 15 50 40 60 45 30	index = 8, maxIndex = 6

Selection Sort: Determining the Position of the Largest Entry

```
int maxIndex = 0;
for (int index = 1; index <= last; index++)
{
    if (A[index] > A[maxIndex])
        maxIndex = index;
}
// maxIndex is position of largest in A[0..last]
```

The Simple Sorting Strategy

Adapted for Selection Sort

```
for (last = N - 1; last >= 1; last --)
{
    //Move the largest entry in A[0...last] to A[last]
    //Determine Position of the Largest entry in A[0..last]
    int maxIndex = Pos. of the Largest entry in A[0..last]
    swap A[maxIndex] with A[last]
}
```

Selection Sort

```
for (last = N - 1; last >= 1; last --)
{
    //Move the largest entry in A[0...last] to A[last]
    //Determine position of largest in A[0..last] and store in maxIndex
    int maxIndex = 0;
    for (int index = 1; index <= last; index++)
    {
        if (A[index] > A[maxIndex])
            maxIndex = index;
    }
    // maxIndex is position of largest in A[0..last]
    // swap A[maxIndex] with A[last]
    int temp = A[maxIndex];
    A[maxIndex] = A[last];
    A[last] = temp;
}
```

Insertion Sort

Insertion Sort

- Note that for any array $A[0..N-1]$, the portion $A[0..0]$ consisting of the single entry $A[0]$ is already sorted.
- Insertion Sort works by extending the length of the sorted portion one step at a time:
 - $A[0]$ is sorted
 - $A[0..1]$ is sorted
 - $A[0..2]$ is sorted
 - $A[0..3]$ is sorted, and so on, until $A[0..N-1]$ is sorted.

Insertion Sort

The strategy for Insertion Sort:

```
//A[0..0] is sorted
for (index = 1; index <= N -1; index ++ )
{
    // A[0..index-1] is sorted
    insert A[index] at the right place in A[0..index]
    // Now A[0..index] is sorted
}
// Now A[0..N -1] is sorted, so entire array is sorted
```


How Insertion Sort Works

15 10 55 35 30 20 index = 1, Insert $A[1] = 10$ into $A[0..1]$:

10 15 55 35 30 20

10 15 55 35 30 20 index = 2, Insert $A[2] = 55$ into $A[0..2]$:

10 15 55 35 30 20

10 15 55 35 30 20 index = 3, Insert $A[3] = 35$ into $A[0..3]$:

10 15 35 55 30 20

10 15 35 55 30 20 index = 4, Insert $A[4] = 30$ into $A[0..4]$:

10 15 30 35 55 20

10 15 30 35 55 20 index = 5, Insert $A[5] = 20$ into $A[0..5]$:

10 15 20 30 35 55

10 15 20 30 35 55 Array is now sorted

A Closer Look at the Logic of the Insertion Step

$A[0..4]$ is already sorted, insert $A[5]$ into $A[0..5]$:

10 15 30 35 55 20 index = 5, Insert $A[5] = 20$ into $A[0..5]$

unsortedValue = 20, will scan for the right place to put it.

Use a variable scan to find the place where $A[\text{scan}-1]$ is less or equal to unsortedValue:

10 15 30 35 55

// scan = 5

10 15 30 35 55

// scan = 4

10 15 30 35 55

// scan = 3

10 15 30 35 55

// scan = 2

Drop in the unsorted value:

10 15 20 30 35 55

// $A[\text{scan}] = \text{unSortedValue}$

Insertion Sort:

insert $A[\text{index}]$ at the right place in $A[0..\text{index}]$

// $A[0..\text{index}]$ is already sorted

```
int unSortedValue = A[index];
```

```
scan = index;
```

```
while (scan > 0 && A[scan-1] > unSortedValue)
```

```
{
```

```
    A[scan] = A[scan-1];
```

```
    scan --;
```

```
}
```

// Drop in the unsorted value

```
A[scan] = unSortedValue;
```

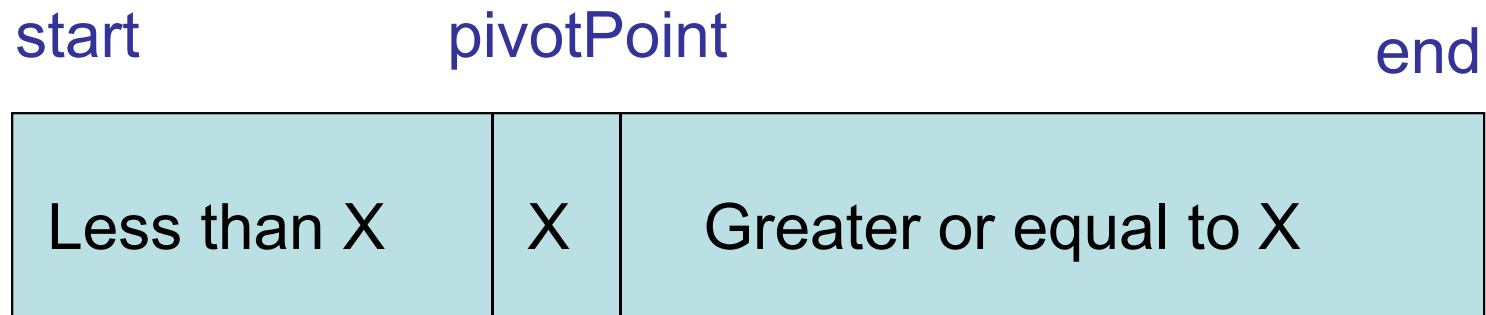
Insertion Sort

```
//A[0..0] is sorted
for (index = 1; index <= N -1; index ++)
{
    // A[0..index-1] is sorted
    // insert A[index] at the right place in A[0..index]
    int unSortedValue = A[index];
    scan = index;
    while (scan > 0 && A[scan-1] > unSortedValue)
    {
        A[scan] = A[scan-1];
        scan --;
    }
    // Drop in the unsorted value
    A[scan] = unSortedValue;
    // Now A[0..index] is sorted
}
// Now A[0..N -1] is sorted, so entire array is sorted
```

Quicksort

Quicksort

To sort a segment $A[\text{start}..\text{end}]$ of an array, Quicksort partitions it into three parts:



Quicksort

For example, in the array

50 67 93 83 90 32 68 13 75 57

You can select $A[0] = 50$ to be the pivot value and then partition as follows:

32 13 50 67 93 83 90 68 75 57

Result of the Partition Step

32 13 50 67 93 83 90 68 75 57

Notice:

1. The pivot value, 50, is in the right place relative to all the other elements! It is where it would be if the array were sorted.
2. The lists on the left and right side of the pivot point are not sorted, but they are shorter!

Quicksort

32 13 50 67 93 83 90 68 75 57

Every time we partition, the pivot value is in the right place relative to all the other elements of the list.

If we recursively carry out the same procedure on the sublists to the left and right of the pivot point, we will keep placing the pivot values in the right position while shortening the remaining sublists. Eventually the sublists get down to a length of 1 or zero, then the whole array is sorted!

Quicksort

- Suppose we have a method to do the partitioning of an array segment `A[start..end]`:
`int partition(int A[], int start, int end)`
- The method partitions the segment and returns the pivot point.

Quicksort

- This recursive procedure will repeatedly partition the sublists until `A[start..end]` is sorted:

```
void doQuicksort(int A[ ], int start, int end)
{
    if (start < end)
    {
        // partition A[start..end] and get the pivot point
        int pivotPoint = partition(A, start, end);
        // recursively do the first sublist
        doQuicksort(A, start, pivotPoint-1);
        // recursively do the second sublist
        doQuicksort(A, pivotPoint+1, end);
    }
}
```

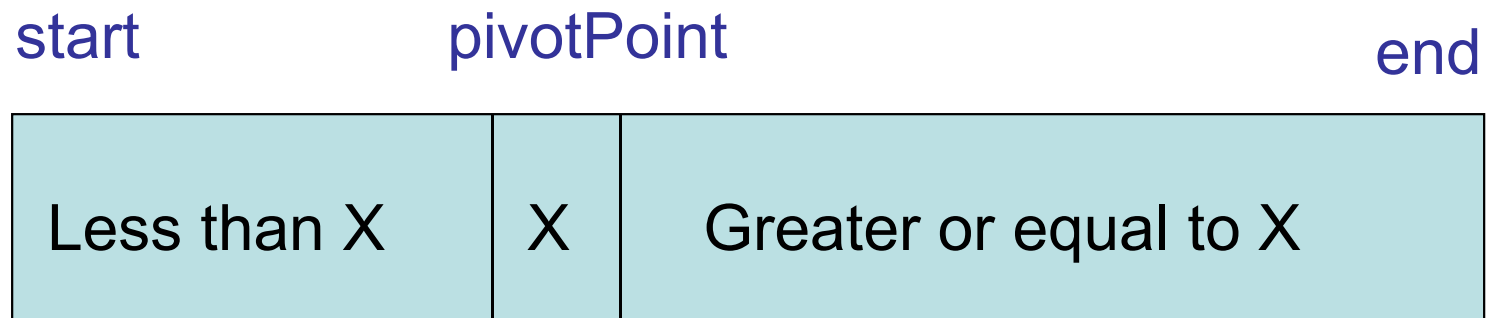
Quicksort

To sort the entire array $A[0..N-1]$, simply call `doQuicksort` and pass it `0` and `N-1` for `start` and `end`:

```
void Quicksort(int A[ ])
{
    doQuicksort(A, 0, N - 1);
}
```

How to Partition

Given an array segment $A[\text{start}..\text{end}]$, we want to partition it and return the pivot point:



How to Partition $A[\text{start}..\text{end}]$

- Arbitrarily choose $X = A[\text{start}]$ as the **pivotValue**, so **start** is the initial **pivotPoint**.
- Use a variable **endOfLeftList** to mark the end of the segment of values smaller than X , and the beginning of the segment of values larger than or equal to X
- Use a variable **scan** to mark the end of the segment that is larger than or equal to X

How to Partition A[start..end]

- A typical stage in the execution of partition.
- Initially: `endOfLeftList = start` and `scan = start+1`.

start	endOfLeftList	scan	end
X	Less than X	Greater or equal to X	Unknown

The End Stage of Partition

At the end, `endOfLeftList` may still be equal to `start`.

If not, then $A[\text{endOfLeftList}] < X$.

Swap $A[\text{start}]$ with $A[\text{endOfList}]$ to get X between the two sublists.

`start`

`endOfLeftList`

`scan`

X	Less than X	Greater or equal to X
---	-------------	-----------------------

How to Partition A[start..end]

```
int partition(int A[ ], int start, int end)
{
    int pivotValue = A[start];
    endOfLeftList = start;
    // At this point A[endOfLeftList] == pivotValue
    for (int scan = start + 1; scan <= end; scan++)
    {
        if (A[scan] < pivotValue)
        {
            endOfLeftList++;
            swap(A, endOfLeftList, scan);
            // At this point A[endOfLeftList] < pivotValue
        }
    }
    // Move the pivotValue between the left and right sublists
    swap(A, start, endOfLeftList);
    return endOfLeftList;
}
```

Array Searching Algorithms

Array Searching Algorithms

Two methods for searching an array for a given item:

1. The **Sequential Search** method can be used with any array.
2. The **Binary Search** method can only be used with arrays that are known to be sorted, but is much faster than Sequential Search.

Sequential Search

- To search an array $A[0..N-1]$ for a value X , start an $index$ at one end of the array, say 0 .
- Step $index$ through the array, examining each $A[index]$ to see if it is equal to X .
- Stop if you find X and return $index$. Otherwise you get to the end of the array and return -1 .

Sequential Search

Search an array $A[0..N-1]$ for X

```
int search(int A[ ], int X)
{
    // Default assumption is X won't be found
    int position = -1;
    boolean found = false;
    int index = 0;
    while (!found && index < N)
    {
        // check A[index]
        if (A[index] == X)
        {
            found = true;
            position = index;
        }
        index ++;
    }
    return position;
}
```

Efficiency of Sequential Search

- In the worst case, you search the entire array, performing N comparisons
- If you are lucky, you find X the first place you look, requiring only one comparison
- On average, you perform $N/2$ comparisons

Binary Search

Binary Search

- Works on a sorted portion $A[\text{lower}..\text{upper}]$:
- Compare X to $A[\text{middle}]$, where middle is the midpoint between lower and upper :

$$\text{middle} = (\text{lower} + \text{upper})/2$$

- If $X == A[\text{middle}]$, return middle (we found it!)
- If $X < A[\text{middle}]$, then continue search in $A[\text{lower}..\text{middle}-1]$
- If $X > A[\text{middle}]$, then continue search in $A[\text{middle}+1..\text{upper}]$
- Search terminates if X is found, or when we try to search an empty segment.

Binary Search of A[lower..upper]

- To continue search in $A[\text{lower}..\text{middle}-1]$, keep lower the same and replace upper with $\text{middle}-1$:
 $\text{upper} = \text{middle} - 1$
- To continue search in $A[\text{middle}+1..\text{upper}]$, replace lower with $\text{middle}+1$ and keep upper the same:
 $\text{lower} = \text{middle} + 1$

Binary Search of A[0..N-1]

```
// returns index of X if found, -1 otherwise
int binSearch(int A[ ], int X)
{
    int lower = 0, upper = N-1;
    int position = -1;      // index of X to be returned
    boolean found = false; // assumption is X will not be found
    // if X is there, it must be in A[lower..upper]
    while (!found && lower <= upper)
    { // if X is there, it must be in A[lower..upper]
        int middle = (lower + upper)/2;
        if (A[middle] == X)
        {
            found = true; position = middle;
        }
        else if (A[middle] > X)
        { // if X is there, it is in A[lower..middle-1]
            upper = middle - 1;
        }
        else
        { // A[middle] < X. if X is there, it is in A[middle+1, upper]
            lower = middle + 1;
        }
    }
    return position;
}
```

Recursive Binary Search

- The logic of binary search has a natural recursive implementation:
- If $\text{lower} > \text{upper}$, then return -1 (base case).
- Compare X to $A[\text{middle}]$, where middle is the midpoint between lower and upper:

$$\text{middle} = (\text{lower} + \text{upper})/2$$

- If $X == A[\text{middle}]$, return middle (we found it!)
- If $X < A[\text{middle}]$, then continue search in $A[\text{lower}..\text{middle}-1]$
- If $X > A[\text{middle}]$, then continue search in $A[\text{middle}+1..\text{upper}]$

Recursive Binary Search of A[lower..upper]

```
int binSearch(int A[ ], int lower, int upper, int X)
{
    // check base case for missing X
    if (lower > upper)
        return -1;
    // check if X is at the middle
    int middle = (lower + upper)/2;
    if (A[middle] == X)
        return middle;
    // determine which segment to continue search in
    if (A[middle] < X)
        return binSearch(A, middle+1, upper, X);
    else
        return binSearch(A, lower, middle-1, X);
}
```

Efficiency of Binary Search

- A basic step in binary search is to split the array, compare X to the middle element, and then select the half of the array in which to continue the search
- Each basic step reduces the size of the array to half the previous size
- If the array has size N , binary search will require no more than $\log N$ basic steps in the worst case

Efficiency of Binary Search

Binary Search is very efficient: large increases in the size of the array require very small increases in the number of basic steps, approximately:

size of array	# steps needed
500	8
1 thousand	10
1 million	20

Analysis of Algorithms

Efficiency of Algorithms

- Usually there is more than one algorithm for solving a given problem.
- One algorithm may be more **efficient** than another, that is, it may need less time, or less memory, to solve a problem of a given size

Criteria for Measuring Efficiency

- **Time**: the time efficiency of an algorithm is measured by the **time complexity function** of the algorithm.
- **Space**: the space efficiency of an algorithm is measured by its **space complexity function**.

Computational Problems

- A **computational problem** is a problem that is meant to be solved by an algorithm.
- A computational problem is described by specifying the data to be input to the algorithm, and the output that should be produced by the algorithm
- Each possible input is called an **instance** of the problem

Instances of Problems

- Each instance is characterized by its size, the amount of memory occupied by the input data that describes the instance.
- The **size** of an instance is the number of bits occupied by the input, but is usually specified by giving an integer that allows us to deduce the actual size in bits.

A Typical Description of a Computational Problem

The problem of summing an array:

INPUT: an array of size N .

SIZE OF INPUT: N is the number of entries in the array.

OUTPUT: an integer representing the sum of the array entries.

The Time Complexity Function

- The time complexity $f(N)$ of an algorithm is determined by counting the number of basic steps executed by the algorithm when solving an instance of size N .
- A **basic step** is an operation that is executed in **constant time**, that is, the time to execute the operation does not increase even if the size of the input increases.
- A basic step is a theoretical approximation for an operation that could be built into the hardware of any reasonable computer.

Basic Operations

- A basic step is also called a **basic operation**.
- Basic steps do not specify the constant time in which they execute: we do not differentiate between a basic step that executes in 1msec and one that executes in 1000 msec.
- Ignoring constant factors in this manner allows the theory to be applicable to computers with different built-in hardware operations and different technologies.
- Ignoring constant factors also means we do not differentiate between 1 basic operation, or 10, or even 100 basic operations.

Computing the Complexity Function of an Algorithm

- We do not need to count all basic operations performed by the algorithm
- For example, if a loop executes N times and each loop iteration executes a constant number of basic operations, the entire loop executes N basic operations
- For most algorithms, we can pick one type of basic operation and count just that to determine the complexity of the algorithm.

Selecting the Basic Operations to Count

- For most algorithms we can count just one or two types of basic operations.
- Select operations that are germane to the problem: for example, sorting and searching algorithms should count comparisons between array entries.
- Select at least one operation in every loop.

Average and Worst Case Complexity

- An algorithm may require a different number of basic steps in solving two different instances of the same size: When searching for X in an array of size N , Sequential search may find X after only 1 comparison, or may require N comparisons.
- The **average case complexity function** averages the number of basic steps required over all instances of size N .
- The **worst case complexity function** is the number of basic steps required for those instances of size N that require the most work to solve.

Average Case Complexity

- Is a good measure to use when you want to know how an algorithm is likely to perform in practice.
- Requires a knowledge of the frequency distribution of problem instances: how often each instance is likely to appear in practice.
- Is difficult to use in practice because reliable estimates of frequency distributions are usually not available

Worst Case Complexity

- Measures the efficiency of an algorithm by how it does on the worst case inputs.
- Is a good measure to use when you want a performance guarantee.

Worst Case Complexity

- The math involved in computing the worst case complexity is easier than the math in average case complexity.
- For these reasons, analysis of algorithms is usually based on worst case complexity.

Comparing Algorithms by Their Complexity Functions

- Let F and G be two algorithms for solving a problem, and let their complexity functions be $f(n)$ and $g(n)$.
- To see the relative performance of the two algorithms, look at the ratio $f(n)/g(n)$ as n gets large.
- We assume that $f(n) > 0$ and $g(n) > 0$ for all $n > 0$.

Comparing Complexity Functions

Simplest case in comparing two algorithms is when the limit $f(n)/g(n)$ exists as n goes to infinity. There are three possible cases where the limit exists:

- Limit of $f(n)/g(n)$ is a positive constant K
- Limit of $f(n)/g(n)$ is infinite
- Limit of $f(n)/g(n)$ is 0

Comparing Complexity Functions

Algorithm F has complexity function $f(n)$

Algorithm G has complexity function $g(n)$

If the limit of $f(n)/g(n)$ is infinite, then Algorithm F is taking a lot more time than G as the size of the problem gets bigger, so G is more efficient.

Comparing Complexity Functions

Algorithm F has complexity function $f(n)$

Algorithm G has complexity function $g(n)$

If limit of $f(n)/g(n)$ is zero, then Algorithm G is taking a lot more time than F as the size of the problem gets bigger, so F is more efficient.

Comparing Complexity Functions

Algorithm F has complexity function $f(n)$

Algorithm G has complexity function $g(n)$

If the limit of $f(n)/g(n)$ is a positive constant K , then Algorithm F is performing K times as many basic operations as G . But constant factors are not significant, so F and G perform the same number of basic operations for really large problems sizes.

The Big O Notation

- In general, the limit $f(n)/g(n)$ may not exist.
- Nevertheless, there may be a positive constant K such that $f(n)/g(n) \leq K$ when n gets large enough.
- If this happens, it means growth of $g(n)$ keeps pace with growth of $f(n)$ and keeps $f(n)/g(n)$ from going to infinity.
- This means that algorithm F is no worse than G , and we say f is in $O(g)$

Space Complexity Functions

- The space complexity function $f(N)$ of an algorithm is a measure of the amount of memory the algorithm requires to solve a problem of size N .
- Space complexity as a measure of efficiency is not often used in theoretical analysis of algorithms, partly because cost of memory keeps declining.